

Stakeholder consultation on the draft Guidelines on the transparency requirements for certain AI systems under Article 50 AI Act

Fields marked with * are mandatory.

Stakeholder consultation on the draft Guidelines on the transparency requirements for certain AI systems under Article 50 AI Act

Disclaimer: This document of the AI Office is prepared for the purpose of consultation and does not prejudice the final decision that the Commission may take on the final text of the guidelines on transparency requirements under Article 50 AI Act. The responses to this consultation will provide input for the guidelines on the transparency requirements for certain AI systems under Article 50 AI Act.

This consultation is targeted to stakeholders of different categories, including, but not limited to, providers and deployers of interactive and generative AI models and systems, providers and deployers of biometric categorisation and emotion recognition systems, private and public sector organisations using such interactive and generative AI systems, as well as academia and research institutions, civil society organisations, governments, supervisory authorities and the general public.

The Artificial Intelligence Act ('the AI Act'), which entered into force on 1 August 2024, creates a single market and harmonised rules for trustworthy and human-centric Artificial Intelligence (AI) in the EU. It aims to promote innovation and the uptake of AI, while ensuring a high level of protection of health, safety and fundamental rights, including democracy and the rule of law.

Among various obligations, trustworthiness of AI systems is ensured through a set of transparency obligations in Article 50 AI Act. These transparency obligations are applicable as of 2 August 2026.

They aim to enable natural persons to recognise interaction with and content generated or manipulated by AI systems, thus reducing the risks of impersonation, deception or anthropomorphisation and fostering trust and integrity in the information ecosystem.

Pursuant to Article 96(1)(d) AI Act, the Commission shall issue guidelines on the practical implementation of transparency obligations laid down in Article 50 AI Act.

The purpose of the present targeted stakeholder consultation is to collect input from a wide range of stakeholders on the draft Commission guidelines on the application of the transparency obligations in Article 50 AI Act.

The drafting of these Guidelines was informed by input from a variety of stakeholders collected during a broad consultation organised by the Commission and input from the Member States in the AI Board. The draft guidelines are now published for additional stakeholder feedback before they are formally adopted by the Commission.

The targeted consultation **is available in English only** and will be **open for 4 weeks starting on 8 May until 3 June 2026**.

All contributions to this consultation may be made publicly available. Therefore, please do not share any personal or confidential information in your contribution. It is your responsibility to avoid personal data and any reference in your contribution itself that would reveal your identity.

The guidelines will be complemented with a Code of Practice that is under development to support the practical and effective implementation of the requirements in Article 50(2) and (4) on marking and labelling of AI generated content.

This consultation accompanies the draft Guidelines on the transparency obligations under Article 50, which aim to support a consistent interpretation and effective implementation of the transparency requirements applicable to certain AI systems and AI-generated or manipulated content across the Union.

This consultation seeks to collect targeted feedback on the draft Guidelines prepared by the Commission based on broad stakeholder input. In particular, the Commission seeks feedback on whether the explanations, concepts and examples provided sufficiently support stakeholders in understanding and complying with the transparency obligations laid down in Article 50 AI Act, including when and how such obligations apply in practice. The draft Guidelines address the transparency obligations under Article 50 AI Act and are structured as follows:

- **Section I** presents an introduction outlining the background, objectives and legal context of the present Guidelines. It recalls the rationale of transparency within the AI Act, including its links to fundamental rights, user awareness and trust, and situates Article 50 within the broader risk-based framework of the AI Act.

- **Section II** provides an overview of the transparency obligations and related horizontal topics. It includes an explanation of the different obligations under Article 50, the actors responsible for their compliance, exclusions from scope (such as purely personal non-professional activities, research and development), and the interplay with other provisions of the AI Act, including prohibited practices and high-risk AI systems and general-purpose AI models and systems.
- **Section III** addresses transparency obligations for AI systems intended to interact directly with natural persons under Article 50(1) AI Act. It sets out the main concepts and scope of application, details the information obligation, and explains the relevant exceptions (including cases of obvious interaction and certain law enforcement uses), as well as the interplay with other Union legal acts.
- **Section IV** provides guidance on the marking and detection of AI-generated or manipulated content under Article 50(2) AI Act. It clarifies the scope of application, including different modalities of synthetic content, and explains the transparency obligation (marking and detection) and the technical requirements (effectiveness, reliability, robustness and interoperability). It also addresses relevant exceptions and the interaction with other Union legal frameworks.
- **Section V** covers transparency obligations for emotion recognition systems and biometric categorisation systems under Article 50(3) AI Act. It outlines the main concepts, scope and applicable obligations, as well as situations falling outside the scope and the interaction with other Union legal acts.
- **Section VI** addresses the labelling of deep fakes and AI-generated or manipulated text published to inform the public on matters of public interest under Article 50(4) AI Act. It provides clarification on the notions of e.g. deep fakes and matters of public interest, the applicable disclosure obligations (including for evidently creative, artistic, satirical, fictional or analogous works or programmes), and the relevant exceptions. It also explains the relationship with other Union legal acts.
- **Section VII** sets out horizontal requirements applicable to the information provided under Article 50(5) AI Act, including general principles ensuring that transparency information is clear, meaningful and accessible to users.

All participants are invited to provide feedback on the particular sections of the draft Guidelines they are interested in. As not all sections may be relevant for all stakeholders, respondents may reply only to the section (s) they would like. Respondents are encouraged to provide explanations and practical cases as part of their responses to support the practical usefulness of the Guidelines. We kindly ask the respondents to specify the exact section and paragraph of the draft Guidelines to which their comments refer.

The feedback collected through this consultation will support the Commission in refining and finalising the Guidelines, with the objective of ensuring that they are clear, comprehensive and practically useful for all stakeholders involved in the development, deployment and supervision of AI systems within the scope of Article 50 of the AI Act.

All contributions to this consultation may be made publicly available.

Therefore, please do not share any personal or confidential information in your contribution (your written feedback). It is your responsibility to avoid personal data and any reference in your written feedback itself that would reveal your identity.

For information on how the Commission processes your personal data please read our Privacy Statement

[7_5_26_Public_Consultation_Transparency50_PrivacyStatement.pdf](#)

Introductory section. Information about the respondent

* First name

Ciaran

* Surname

O'Riordan

* Email address

ciaran.oriordan@eclipse-foundation.org

* Do you agree that we may publish your identity together with your contribution in the instance that all contributions are made publicly available?

If you act in your personal capacity: All contributions to this consultation may be made publicly available. You can choose whether you would like your details to be made public or to remain anonymous. The respondent category that you selected for this consultation, your answer regarding residence, and your contribution may be published as received. Should you choose to remain anonymous, your name will not be published. Please do not include any personal data in the contribution itself.

If you represent one or more organisations: All contributions to this consultation may be made publicly available. You can choose whether you would like respondent details to be made public or to remain anonymous. Only the following organisation details may be published: The respondent category that you selected for this consultation, the name of the organisation on whose behalf you reply as well as its size, its presence in or outside the EU and your

contribution as received. Should you choose to remain anonymous, your name will not be published. Please do not include any personal data in the contribution itself if you want to remain anonymous.

- ☒ Yes
☐ No

* Do you agree that we may contact you in the event of follow-up questions or if we want to learn more about your responses?

- ☒ Yes
☐ No

* Do you represent an organisation (e.g., a public organisation, a company, a think tank or a civil society /consumer organisation) or act in your personal capacity (e.g., independent expert)?

- ☒ Organisation
☐ In a personal capacity

* If you are representing an organisation, please specify the name of the organisation:

Open Regulatory Compliance Working Group (ORC) (hosted by Eclipse Foundation AISBL)

* Type of organisation

Civil society organisation/association

* Is a representation of the organisation located in the EU?

- ☒ The organisation's headquarter is located in the EU
☐ A branch office, or any representation of the organisation is located in the EU
☐ None of the representations of the organisation is located in the EU

* Select the EU Member State where the organisation's headquarter, or representation is located

BE - Belgium

* Select the size of the organisation

Medium (50-249 employees)

* Sector(s) of activity

- | | | |
|--|---|--|
| ☒ Information technology | ☐ Employment | ☐ Transport |
| ☐ Public administration | ☐ Education and training | ☐ Telecommunications |
| ☐ Law enforcement | ☐ Consumer services | ☐ Retail |
| ☐ Justice sector | ☐ Business services | ☐ E-commerce |
| ☐ Legal services sector | ☐ Banking and finances | ☐ Advertising |
| ☐ Cultural and creative sector, including media | ☐ Manufacturing | ☐ Consumer protection |
| ☐ Healthcare | ☐ Energy | ☐ Others |

* Describe the activities of your organisation or yourself

1300 character(s) maximum

The Open Regulatory Compliance Working Group (ORC) consists of key stakeholders from industry and open source communities collaborating to support compliance with government regulations, ensuring the continued use and advancement of open source through the software supply chain.

When providing feedback, please ensure that your responses are concise and clearly reference the relevant section, subsection and/or paragraph number(s) of the draft Guidelines.

This will support the Commission in efficiently analysing stakeholder input and refining the Guidelines to ensure they are clear, comprehensive and practically usable.

* **On which part(s) of the draft Guidelines would like comment?** *Multiple answers are possible. Please note that selecting a particular answer will direct you to a set of questions specifically related to subject specified.*

- ☐ **Section I - Background and objectives**
- ☒ **Section II - Overview of transparency obligations and horizontal topics**
- ☒ **Section III - Article 50(1): Transparency for interactive AI systems**
- ☒ **Section IV - Article 50(2): Marking and detection of AI-generated or manipulated content**
- ☐ **Section V - Article 50(3): Emotion recognition and biometric categorisation systems**
- ☐ **Section VI - Article 50(4): Labelling of deep fakes and certain text publications**
- ☒ **Section VII - Horizontal requirements (Article 50(5))**

Section II - Overview of transparency obligations and horizontal topics

Question A. Are there any aspects / sections / paragraphs that require further clarification or refinement?

- ☒ Yes
- ☐ No

Please indicate specific parts.

2000 character(s) maximum

The research / placing-on-market boundary (paragraphs 19-21)

The Guidelines confirm that the research and development exclusions (Article 2(6) and 2(8); paragraphs 19-20) cease to apply once a system is placed on the market or put into service, that the purely-personal-use exclusion (paragraphs 16-18) lifts only the deployer's duties and not the provider's, and that there is no open-source exemption (paragraph 21). The combined effect is that there is no de minimis threshold at the provider level for non-commercial or research-driven open release: publishing open weights or an open system to the public could constitute placing on the market and triggers the full set of provider obligations, regardless of whether any economic benefit is sought.

Iterative open development does not map neatly onto a single "placing on the market" moment. We ask the Commission to clarify how the boundary between research and testing in a pre-market phase and "placing on the market" / "putting into service" should be applied to common open-development practices, for example, the release of research previews, intermediate checkpoints, and the distinction between releasing a model and releasing a system, so that non-commercial and research-driven open release is not unintentionally chilled. Distinguishing characteristics could be the absence of production support commitments, the experimental or unstable nature of checkpoints, explicit research-only disclaimers, and whether the release is intended primarily to facilitate scientific evaluation and reproducibility rather than downstream commercial deployment.

Section III - Article 50(1): Transparency for interactive AI systems

Question A. Are there aspects requiring further clarification for effective implementation and compliance?

- ☒ Yes
☐ No

Please indicate specific parts.

2000 character(s) maximum

"Lifecycle of the AI system" (paragraph 30, with paragraphs 10, 13 and 70)

Paragraph 30 says that providers are responsible for ensuring that natural persons are effectively informed "throughout the lifecycle of the AI system, including after its placement on the market and its putting into service." We would welcome a more precise definition or interpretation of the scope of "lifecycle of the AI system" in this context. In particular, how should this responsibility be interpreted for FOSS or open-weight providers, where downstream modified variants may evolve outside the control of the original provider?

We note that the Act provides no Article 50-specific mechanism to resolve this. The value-chain reclassification rule in Article 25, and the definition of "substantial modification" in Article 3(23), are confined to high-risk AI systems and are keyed to the requirements of Chapter III, Section 2; they therefore do not determine when responsibility passes for a transparency-only system that is forked or fine-tuned downstream. In the absence of such a rule, the position falls back on the general provider definition in Article 3(3).

We also note that "substantial modification" is used in the AI Act and the Cyber Resilience Act, and that interpretation in that context is subject to ongoing discussion. Ideally, the term should have a consistent meaning between the two instruments, but whether this really makes sense is down to whether the same definition can function correctly in both contexts. Since interpretation in both contexts is still in flux, we can only advise to take care when deciding how to harmonise, or to not harmonise.

Question B. Are there additional practical examples or use-cases that should be clarified?

- ☒ Yes
☐ No

Please specify:

2000 character(s) maximum

We ask the Commission to clarify that:

1. "Lifecycle of the AI system" in paragraph 30 refers to the lifecycle of the system as placed on the market or put into service by that provider; and

2. Where a downstream actor substantially modifies an open system and places that variant on the market or puts it into service under its own responsibility and name, that actor becomes the provider of the variant under Article 3(3) and assumes the Article 50 obligations in respect of it, so that the original provider's lifecycle responsibility does not extend to downstream variants outside its control.

This reading is consistent with paragraph 70, which already accepts that a provider may rely on a marking solution implemented upstream while remaining responsible for demonstrating compliance for its own system. Making the handover explicit for the open source modification chain would close a gap the Guidelines presently leave open.

Section IV - Article 50(2): Marking and detection of AI-generated or manipulated content

Question A. Are there aspects requiring further clarification or refinement to ensure effective implementation of marking and detection obligations?

- ☒ Yes
☐ No

Please specify:

2000 character(s) maximum

Interoperability and international context (paragraphs 74 and 77)

Paragraph 74's interoperability requirement, regarding the capability of marking and detection solutions to operate seamlessly across multiple systems, actors and contexts, together with paragraph 75's reference to the state of the art "as may be reflected in relevant technical standards," creates an opportunity to align Article 50 compliance with open, royalty-free, internationally-developed provenance and marking standards.

Interpreting these provisions to favour such open standards would serve three EU objectives simultaneously: legal compliance, the continued viability of open source development, and convergence with partner jurisdictions with which the EU cooperates on AI standards. Conversely, an interpretation that privileged proprietary or EU-specific marking schemes would disadvantage open source developers and risk fragmenting the very standards landscape that is a core part of the EU's international cooperation, such as the AI Factories network for partner-hosted compute.

We ask the Commission to confirm that the interoperability requirement of Article 50(2) can be satisfied through adherence to recognised open international standards developed in inclusive, multi-stakeholder forums, and to avoid any interpretation that would effectively mandate solutions available only to the largest providers.

Robustness and technical feasibility (paragraphs 74-76)

The robustness requirement in paragraph 74, which expressly covers "adversarial attacks," could be an issue for FOSS and open-weight systems. Once weights are openly distributed, a downstream actor can strip or retrain out any embedded mark, the limiting case of an adversarial attack, and one that cannot be prevented.

Question B. Are there additional examples, use-cases or technical approaches that should be included?

- ☒ Yes
☐ No

Please specify:

2000 character(s) maximum

Paragraph 74, together with paragraphs 75 and 76, could be read as an implicit recognition of this situation, through the introduction of the notion of “technically feasible” and the reference to feasibility being assessed “within the specific technical architecture and operational environment concerned.” However, paragraph 76 also states that technical feasibility is “an objective notion that is not dependent on the specific resources and capabilities of individual providers.” This leaves ambiguity as to whether the open-weight release model is to be treated as part of the objective “technical architecture and operational environment” (which paragraph 76 permits to be taken into account) or as a feature of an “individual provider’s capabilities” (which paragraph 76 excludes from consideration). It would be useful to clarify how robustness should be interpreted in the context of a FOSS AI system:

1. The openly-distributed / open-weight release model is an objective architectural characteristic of the system, not a provider-specific resource or capability limitation, so that robustness against adversarial removal is assessed by reference to what is achievable for openly-distributed models; and

2. An upstream provider that implements required marking and detection at the point of release does not, by that fact alone, infringe Article 50(2) where a downstream actor with access to the weights subsequently removes or defeats the mark.

We would also caution against any implicit assumption that “structural” marking encoded in the weights during training (as contemplated in the parallel Code of Practice) resolves the robustness question for open-weight systems. Such marking does not reliably survive adversarial fine-tuning, and treating it as a solution would set an expectation that open source developers cannot meet.

This does not alter the provider’s obligations regarding marking measures at the point of release. It acknowledges objective characteristics.

Question C. What technical approaches or solutions should be considered for the implementation of the marking and detection obligations under Article 50(2) AI Act, in particular in relation to AI agents and virtual or immersive environments (e.g. virtual reality)? Please provide concrete examples, methodologies or tools where possible.

1500 character(s) maximum

Section VII - Horizontal requirements (Article 50(5))

Question A. Are there aspects requiring further clarification or refinement to ensure effective implementation? Please indicate specific parts.

2000 character(s) maximum

Proportionate enforcement without the Code of Practice (paragraphs 135-136)

Paragraphs 135-136 establish a two-track approach in which signatories to a Code of Practice benefit from focused supervision while non-signatories are expected to demonstrate compliance by other means. The draft notes that this may include performing a gap analysis against the code and could result in a larger number of information requests from authorities.

We recognise that the Code of Practice has been drafted with input from open-source experts, and that it is well-suited to certain organisational setups. For open-source entities that operate with centralised legal and administrative structures, the Code provides a valuable way of demonstrating compliance.

There are also many FOSS projects and smaller open-source companies with different operational setups which may include decentralised communities and volunteer maintainer networks. The Code of Practice is not intended for every FOSS structure, including those who do not have the legal or procedural systems to formally sign, administer, and audit adherence to the Code of Practice.

We ask the Commission to confirm that the route to demonstrating compliance by “adequate alternative means” remains genuinely accessible to them. Specifically, we request clarification that this route accommodates transparent compliance practices native to open source development such as comprehensive public model cards, open documentation, source code disclosure and public version-controlled repositories. For these developers, non-adherence to the Code of Practice should not operate as a de facto penalty or automatically trigger burdensome corporate-style gap analyses that are structurally impractical for them to produce.

A proportionate supervisory approach should therefore take account of the organisational realities of distributed open-source development and avoid indirectly privileging large vertically-integrated providers.

Contact

[Contact Form](#)